

# Midterm 1: Virtualization

Time: 75 minutes Total Questions/Total points: 70

## Instructions

This exam is a closed book, closed notes. You will hand the question set and the scantron sheet back to the instructor after the exam is over.

Please fill in your name and student ID on the answer sheet.

When you begin the exam, you will need to agree with the following statements:

1. 1. I understand that this exam must be worked on independently.
2. 2. I will not communicate or collaborate with any others during the exam in any way;

for example:

- I will not chat, text, or post any questions or information.
  - I will not ask for, receive, or provide help to others.
  - I will notify the instructor if I am aware of others violating this policy.
3. I will not discuss the questions with students of the other section.

Unless stated (or implied) otherwise, you should make the following assumptions:

- The OS manages a single uniprocessor (single core).
- All memory is byte addressable.
- The terminology lg means  $\log_2$ .
- 1 byte = 8 bits.
- 1 hexadecimal digit takes 4 bits.
- Page table entries require 4 bytes unless otherwise stated.
- Data is allocated with optimal alignment, starting at the beginning of a page.
- Leading zeros can be removed from numbers (e.g., 0x06 == 0x6).
- Hex numbers are represented with a preceding "0x"; numbers without 0x are decimal.

## Table of powers of 2

Power of 2	Bytes	KB	MB
0	1		
1	2		
2	4		
3	8		
4	16		
5	32		
6	64		
7	128		
8	256		
9	512		
10	1,024	1	
11	2,048	2	
12	4,096	4	
13	8,192	8	
14	16,384	16	
15	32,768	32	
16	65,536	64	
17	131,072	128	
18	262,144	256	
19	524,288	512	
20	1,048,576	1,024	1
21	2,097,152	2,048	2
22	4,194,304	4,096	4
23	8,388,608	8,192	8
24	16,777,216	16,384	16
25	33,554,432	32,768	32
26	67,108,864	65,536	64
27	134,217,728	131,072	128
28	268,435,456	262,144	256
29	536,870,912	524,288	512
30	1,073,741,824	1,048,576	1,024
31	2,147,483,648	2,097,152	2,048
32	4,294,967,296	4,194,304	4,096

## Part 1: CPU virtualization short answers

1. Which of these are not application benefits of an operating system?
  - a. A set of simpler abstractions against which to program
  - b. Independence from specific hardware and devices
  - c. More control over how hardware is used
2. Which of these are user benefits of an operating system?
  - a. Protecting running programs against the misbehavior of other programming
  - b. More physical memory available to programs
  - c. Sharing hardware by running multiple programs at the same time
  - d. A & C
  - e. All the above
3. Which of the following is not a valid process state
  - a. Blocked
  - b. Ready
  - c. Preempting
  - d. Waiting
4. Processes are a stream of execution and execution state
  - a. True
  - b. False
5. The state of a process is its execution state (registers) and memory state
  - a. True
  - b. False
6. Which of the following are present in processes but not programs
  - a. Stack & registers
  - b. Heap & data
  - c. Code & registers
  - d. Data & code
7. With direct execution, the operating system can take control of the CPU at any time
  - a. True
  - b. False

8. It is not safe to let user-mode code enable and disable interrupts
  - a. True
  - b. False
9. What does a processor do if a process executes a privileged instruction in user mode?
  - a. The processor traps
  - b. The instruction is ignored
  - c. The computer reboots
10. Which of the following are operations allowed in user-mode execution?
  - a. Access to all areas of physical memory
  - b. Saving/restoring registers
  - c. Issuing input/output commands
  - d. Modifying processor memory management state
11. On a timer interrupt, the hardware is responsible for transitioning from user to kernel mode.
  - a. True
  - b. False
12. Processes can move from the running state directly to the blocked state
  - a. True
  - b. False
13. Processes can't move from the blocked state directly to the ready state
  - a. True
  - b. False
14. Only one process can be in the RUNNING state
  - a. True
  - b. False
15. With cooperating multitasking, context switches only happen when a program calls yield()
  - a. True
  - b. False

16. On a timer interrupt, the operating system always context switches to a new process
- a. True
  - b. False
17. On a context switch, user-mode registers are saved in the process-control block (PCB)
- a. True
  - b. False
18. You are running a mix of a word processor and a scientific calculation. Which policy do you prefer?
- a. Round robin
  - b. FIFO/FCFS
  - c. SJF
  - d. SCTF
19. You are designing a word processor and want to make sure it responds quickly when a key is struck. Which scheduling goal are you most interested in?
- a. Turnaround time
  - b. Response time
  - c. Fairness
  - d. No starvation
20. SCTF suffers from head-of-line blocking
- a. True
  - b. False
21. SCTF doesn't suffer from starvation
- a. True
  - b. False
22. With FIFO scheduling, when a task wakes up from I/O it keeps its old place in the ready queue
- a. True
  - b. False

23. With MLFQ, after running tasks in one queue, the scheduler starts running tasks in the next higher-priority queue

- a. True
- b. False

24. With MLFQ, tasks with the same, short burst length will run repeatedly in the same queue using round-robin.

- a. True
- b. False

25. With MLFQ, Boost gives a running process an extra-long time slice

- a. True
- b. False

26. Which of these is a difference between MLFQ without boost and SCTF?

- a. SCTF prioritizes jobs nearly done with their CPU burst
- b. MLFQ is non-preemptive
- c. MLFQ does not suffer from starvation
- d. MLFQ prioritizes shorter jobs

27. Suppose you have a process A running. It then does this sequence of operations

`Calls fork to create process B()`

`Process A issues an I/O request to the disk`

`The system then continues for a little longer`

What states are Process A and Process B in?

- a. A: Ready and B: Running
- b. A: Blocked and B: Ready
- c. A: Blocked and B: Running
- d. A: Running and B: Blocked

28. On a multi-core processor with a single global scheduling queue, the core a processor runs on is chosen

- a. When the process is created
- b. When the process runs the first time
- c. Every time the process runs

29. You are running a set of short jobs (CPU burst is about 1 time slice). Which of these schedulers would give you the best response time?

- a. Single global ready queue
- b. Per-CPU ready queues
- c. Per-CPU ready queues with migration

## Part 2: CPU Virtualization calculations

30. For this schedule using FIFO, what is the average turnaround time?

- a. Less than 3 ms
- b. Between 3-4 ms
- c. Between 5-6 ms
- d. Between 11-12 ms

Job	Length	Arrival
1	5 ms	12 ms
2	4 ms	10 ms
3	2ms	11 ms

31. For this scheduling using SJF, what is the average response time?

- a. Between 2-3ms
- b. Between 5-6ms
- c. Between 7-8ms
- d. Between 8-9ms

Job	Length	Arrival
1	2 ms	0 ms
2	7 ms	0 ms
3	4 ms	0 ms

32. For this scheduling using SCTF, when does job 3 finish? If two jobs have the same remaining time, run the one that started first

- a. At 4 ms
- b. At 9 ms
- c. At 13 ms
- d. At 21 ms

Job	Length	Arrival
1	2 ms	3 ms
2	7 ms	0 ms
3	4 ms	5 ms

33. Using Round Robin scheduling with a 2ms time slice (quantum), when does job 2 finish?

- a. At 7ms
- b. At 8ms
- c. At 12ms
- d. At 14ms

Job	Length	Arrival
1	2 ms	1 ms
2	7 ms	0 ms
3	3 ms	5 ms

34. Assume a parent process always continues running before a child process runs, and children run in order created. Parent has PID 1, child PIDs start at 2 and are allocated sequentially. For this code:

```
int child1_pid = fork();  
int child2_pid = fork();  
printf("c1 = %d, c2 = %d\n", child1_pid, child2_pid );
```

What is printed first?

- a. c1 = 2, c2 = 3
  - b. c1 = 0, c2 = 4
  - c. c1 = 2, c2 = 0
  - d. c1 = 0, c2 = 0
35. For the code in the previous question, what is printed second?
- a. c1 = 2, c2 = 3
  - b. c1 = 0, c2 = 4
  - c. c1 = 2, c2 = 0
  - d. c1 = 0, c2 = 0



## Part 3: Memory Virtualization short answers

36. For a single sequential pass over data, a 4-entry TLB performs better than a 1-entry TLB
- a. True
  - b. False
37. To translate a virtual address using paging, the MMU adds the offset to the physical page number from the TLB
- a. True
  - b. False
38. With multi-level page tables, larger pages lead to fewer levels in the page table
- a. True
  - b. False
39. A process can't access the data of another process because the other process's addresses are marked invalidated in the page table.
- a. True
  - b. False
40. Which mechanism provides the least protection?
- a. Static relocation
  - b. Base and bounds
  - c. Segments
  - d. Paging
41. Which mechanism has the least wasted memory for a standard process layout?
- a. Segments
  - b. Base and bounds
  - c. Paging with linear page tables
42. Which mechanisms support sharing?
- a. Segments
  - b. Paging
  - c. Paging and base/bounds
  - d. Base/bounds and segments
  - e. Segments and paging

43. What is the major benefit of paging over segmentation?
- a. Less wasted memory
  - b. Faster performance
  - c. Less external fragmentation
  - d. Less MMU state
44. With segmentation, the address space of each process must be allocated contiguously in physical memory
- a. True
  - b. False
45. Compared to using segments, using a 2-level page table with no TLB doubles the number of memory references
- a. True
  - b. False
46. When the dirty bit is set in a PTE, it means the contents of a TLB entry do not match the page table
- a. True
  - b. False
47. When using paging, larger pages causes more external fragmentation
- a. True
  - b. False
48. To use huge pages (such as 2mb pages), a programmer always has to request them when allocating memory
- a. True
  - b. False
49. With shared memory, processes map the same physical pages to the same virtual addresses in multiple processes
- a. True
  - b. False
50. FIFO page replacement always has more page faults than LRU page replacement
- a. True
  - b. False

51. An LRU TLB replacement policy is better than a random TLB replacement policy for sequential workloads.
- a. True
  - b. False
52. With a software-filled TLB, the OS specifies the layout of the page table
- a. True
  - b. False
53. With ASIDs, there is no performance penalty for context switching when using a TLB
- a. True
  - b. False
54. When the present bit is cleared in a PTE, accessing the associated page will cause a segmentation fault
- a. True
  - b. False
55. Demand paging never loads the data at an address referenced by a process before it is used the first time
- a. True
  - b. False
56. When a system is thrashing, the CPU is fully utilized
- a. True
  - b. False
57. With inverted page tables, every physical page in a system has a PTE
- a. True
  - b. False
58. Copy-on-write is faster than regular copy when a process reads from most of the copied pages.
- a. True
  - b. False

59. If the clock hand in clock page replacement is moving quickly, it could mean
- a. The program has high locality
  - b. The program has low locality
60. With clock page replacement, a page that is marked dirty always has the referenced bit set
- a. True
  - b. False

## Part 3: Memory Virtualization calculations

61. Given a 4 entry TLB and 4KB pages, what will the TLB miss rate be for the line of code shown in boldface?

- a. Between 0 and 1%
- b. Between 1 and 5 %
- c. Between 5 and 10%
- d. Between 10 and 25%
- e. Higher than 25%

```
// doubles are 8 bytes
double array[1024*1024];
double sum = 0;
for (i = 0; i < 1024; i++)
    for (j = 0; j < 1024; j++)
        Sum = sum + array[1024*i+j];
```

62. Consider this variation of the code. Given a 4 entry TLB and 4KB pages, what will the TLB miss rate be for the line of code shown in boldface?

- a. Between 0 and 1%
- b. Between 1 and 5 %
- c. Between 5 and 10%
- d. Between 10 and 25%
- e. Higher than 25%

```
// doubles are 8 bytes
double array[1024*1024];
double sum = 0;
for (i = 0; i < 1024; i++)
    for (j = 0; j < 1024; j++)
        Sum = sum + array[i+1024*j];
```

63. Using the same variation of the code from the previous question, consider a system with a 16 entry TLB and 2MB pages. What is the TLB miss rate for the line of code in boldface when running this loop many times?

- a. Between 0 and 1 %
- b. Between 1 and 5 %
- c. Between 5 and 10%
- d. Between 10 and 25%
- e. Higher than 25%

64. Given a virtual address size of 32 bits and a page size of 256 bytes and a PTE size of 4 bytes, what is the size of a linear page table?

- a. 4MB
- b. 16MB
- c. 64MB
- d. 256MB

65. Given a virtual address size of 32 bits and a page size of 256 bytes and a PTE size of 4 bytes, how many levels of page table are needed?

- a. 1
- b. 2
- c. 3
- d. 4

66. Given a virtual address size of 32 bits and a page size of 256 bytes and a PTE size of 4 bytes, what is the smallest number of pages needed for page tables for a process with a standard layout and 1kb of code, 1 kb of data, 1kb of heap, and 1kb of stack

- a. 1
- b. 3
- c. 5
- d. 7

67. In the code shown below, where is `longLine` stored?

- a. Static data
- b. Heap
- c. Stack
- d. Code

68. In the code shown below, where is `maxlen` stored?

- a. Static data
- b. Heap
- c. Stack
- d. Code

```
int maxlen = 512;

int main(int argc, char **argv) {
    char shortLine[10];
    char * longLine = malloc(maxlen);
    parseArgs(argc, argv, shortLine, longLine);
}
```

69. With 3 pages of memory, demand paging, FIFO replacement, and this reference string, what pages will be in memory at the end? 5, 4, 2, 1, 3, 2, 4, 1, 2, 3?

- a. 1, 3, 4
- b. 2, 3, 4
- c. 2, 1, 3
- d. 4, 2, 1

70. With 3 pages of memory, demand paging LRU replacement and this reference string, assuming memory starts empty, how many page faults will there be? 3, 4, 2, 1, 0, 3, 2, 3

- a. 3
- b. 4
- c. 5
- d. 6
- e. 7